

Vessa — conversational AI for patient communication

The Problem

A typical US outpatient clinic loses 18% of booked appointments to no-shows and late cancellations. Each empty slot costs roughly \$200 in unbilled time. The standard response is a reminder text, which recovers a few points and irritates everyone else.

The deeper problem is that patients have questions before they attend — about cost, preparation, insurance, whether they still need the visit at all — and there is nobody to answer them. Front desk staff spend 40% of their time on the phone repeating the same fifteen answers.

Our Solution

Vessa answers patients in natural language across SMS, web chat and voice, using the clinic's own scheduling and billing systems as the source of truth.

A patient asks whether they need to fast before their blood test. Vessa reads the order, gives the specific instruction for that panel, and confirms the appointment in the same exchange. When the question needs a human — clinical advice, anything about medication changes — it hands off with a summary rather than a transcript.

Everything is grounded in clinic records. The system will say it does not know rather than guess, and every answer carries a citation back to the record that produced it.

The Market

US outpatient clinics spend about \$12B a year on patient communication, scheduling and front-desk labour. Our segment is independent multi-site groups between 5 and 60 providers — roughly 31,000 organisations — which are large enough to feel the cost and small enough to move without a committee.

Competition

Three types. Legacy patient-engagement suites bundled with the electronic health record, which are broad and poor at conversation. Point-solution reminder tools, which do one thing adequately. And a handful of AI-first entrants funded in the last two years.

We are not alone in this category and we do not claim to be. Our position rests on grounding quality — we integrate at the order level rather than the appointment level, which is what lets us answer preparation questions correctly — and on being the only entrant certified against three major EHR platforms rather than one.

The Team

Dana Okafor ran clinical operations for a 40-site primary care group and spent four years watching front-desk workflows fail in person. Luis Ferreira built conversational systems at a customer-service platform, most recently leading the team responsible for grounding and retrieval.

Nineteen people, including two former clinic managers in customer success.

Neither founder has built a company past this stage. We are hiring a VP of Engineering because we know the gap is real.

Traction

- 61 clinic organisations live, covering 840 providers
- no-show rate across deployed sites fell from 17.4% to 11.9% over six months
- 74% of inbound patient messages resolved without human involvement
- gross revenue retention 94%, net 121%
- average deployment time 3 weeks

Business Model

\$340 per provider per month, billed annually. At 840 providers that is \$3.4M in annual recurring revenue. Gross margin is 68% after inference and telephony costs, up from 54% a year ago as we moved most traffic to smaller models.

Expansion comes from provider count within existing groups rather than upsell, which keeps the motion simple.

Go-To-Market

Direct sales to practice administrators, 7-week average cycle, \$9,400 blended acquisition cost per organisation recovered in 8 months. Two EHR vendors list us in their marketplaces, which now generates a quarter of qualified pipeline.

Funding

We are raising \$16M to add two EHR integrations, build the voice product beyond its current pilot, and expand the clinical review team that audits our grounding quality.